

# Food-Delivery Aggregator - Requirements Pack

---

## Submitted By:

**Name:** Manjot Kaur

**Roll Number:** 08

**Registration Number:** 12412384

**Course:** B.Tech (Generative AI)

**Semester:** 3

**Session:** 2024–2025

**College:** Lovely Professional University

---

## Overview

Build a marketplace that ingests menus, takes orders, assigns drivers, and streams status updates. Key goals: availability (24/7), real-time ETA updates, scalability across geographies, and clear domain boundaries.

---

## Audience

Engineering leads / architects (primary), Product managers (secondary).

---

## 1) Stakeholder analysis & prioritization

Stakeholder	Interests	Priority
Customers	Fast ordering, accurate ETAs, reliable deliveries	High
Restaurants	Accurate menu/catalog, timely orders, commission clarity	High

Couriers	Fair assignments, clear routing, earnings visibility	High
Ops / Support	Monitoring, incident response, capacity control	Medium
Payments	Fraud prevention, settlement speed	Medium
Legal / Compliance	Data residency, privacy, terms	High

**Key takeaways:** prioritize customer UX, restaurant onboarding flows, courier reliability, and operational observability.

---

## 2) Requirements

### Functional

- Menu ingestion & caching (restaurant, categories, prices, availability)
- Cart & checkout (taxes, tips, promos)
- Order lifecycle: placed → accepted (restaurant) → prepared → assigned (courier) → pickup → enroute → delivered → closed
- Dispatch / matching: assign nearest/available courier using policy (latency, ETA, fairness)
- Live tracking: courier geolocation streaming to customer & restaurant
- Notifications: push, SMS, email for status changes
- Payments: charge customer, hold/settle to restaurant, fees
- Admin: dashboards, overrides, restaurant onboarding, courier management
- Reporting & auditing (orders, payouts, disputes)

### Non-functional

- Availability: 99.9% (regionally) during operating hours; 99.95% critical paths

- Real-time performance: p95 location update propagation < 1s; p95 order placement < 300ms
- Scale targets: support >10k concurrent active deliveries; burst handling 5–10×
- Security: PCI/DSS for payments, GDPR for personal data

## Technical

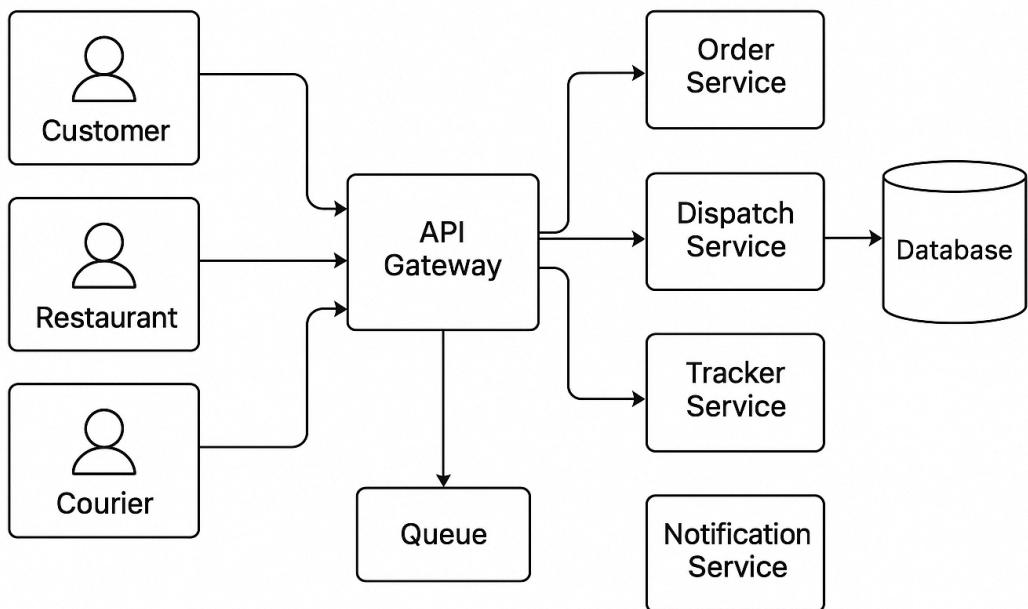
- Real-time channels: WebSocket or Server-Sent Events for location/status
  - Use third-party Maps (routing/ETA) with caching of route results
  - Geo-sharding for regional scaling and data residency
  - Event-driven (message bus) for decoupled services
- 

## 3) Constraints & assumptions

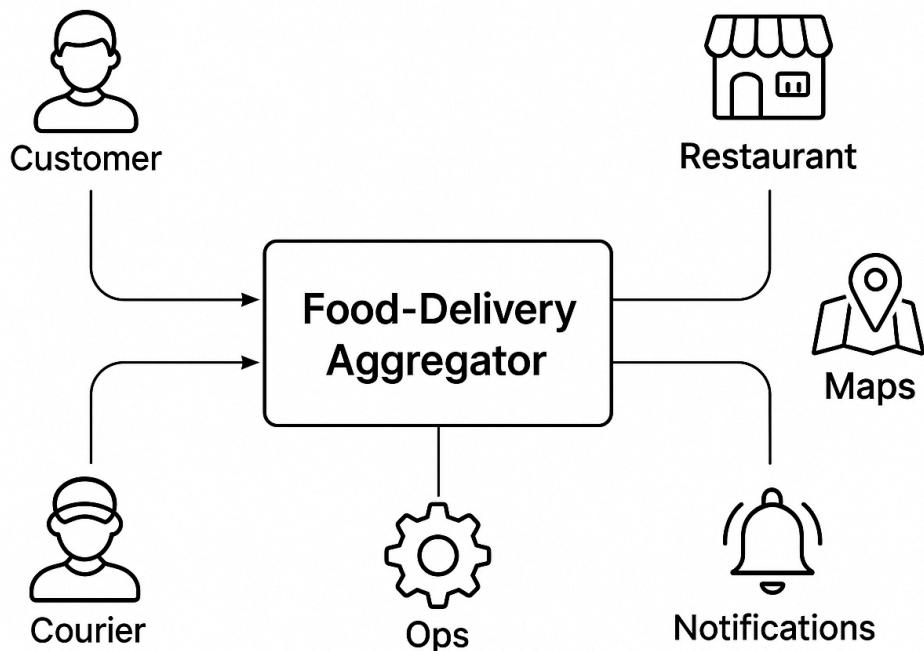
- Use third-party maps/routing APIs (e.g., Google Maps, Mapbox)
  - Eventual consistency acceptable for ETAs and route recalculations
  - Courier devices have intermittent connectivity
  - Restaurants own menu truth but system caches for performance
  - Payment provider SLA  $\geq$  99.9%
- 

## 4) High-level architecture

## HIGH-LEVEL ARCHITECTURE

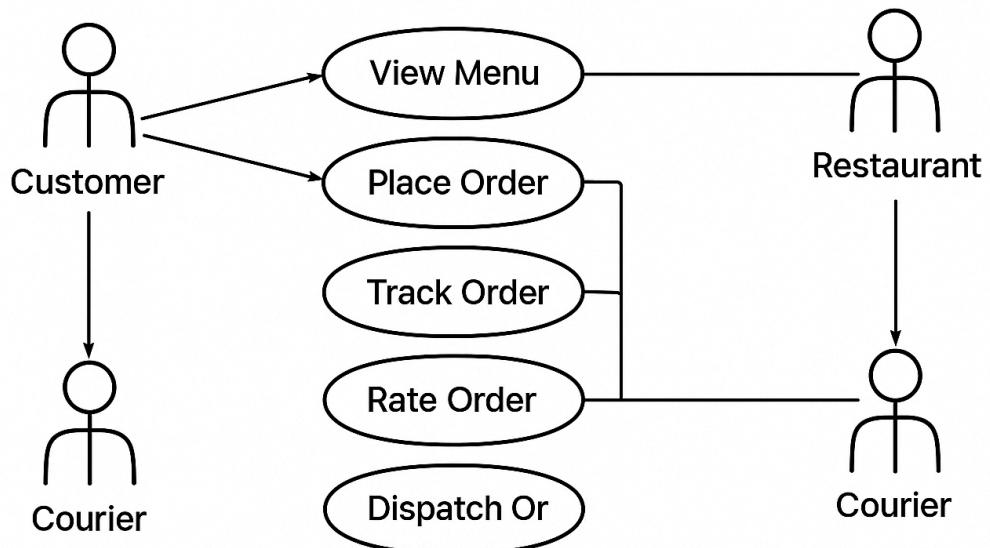


## 5) System Context



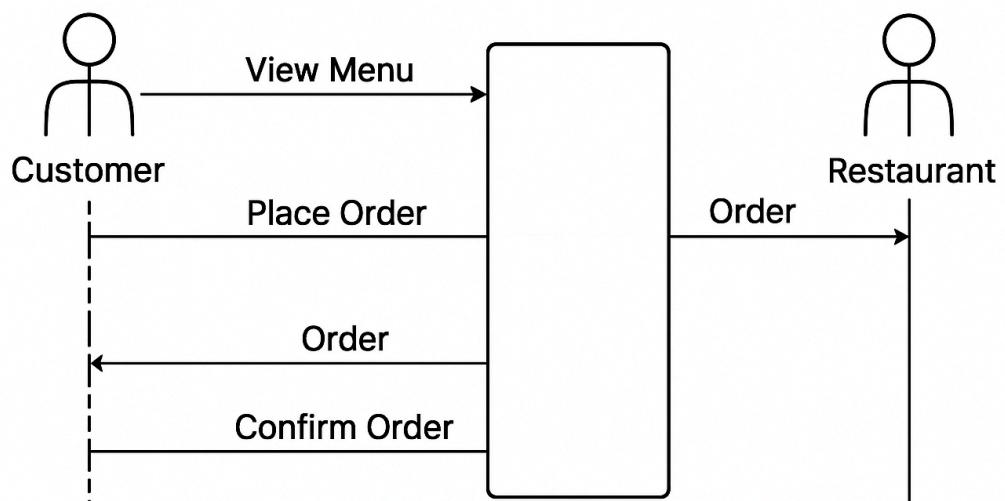
## 6) Use Case Daigram

Use-Case Diagram

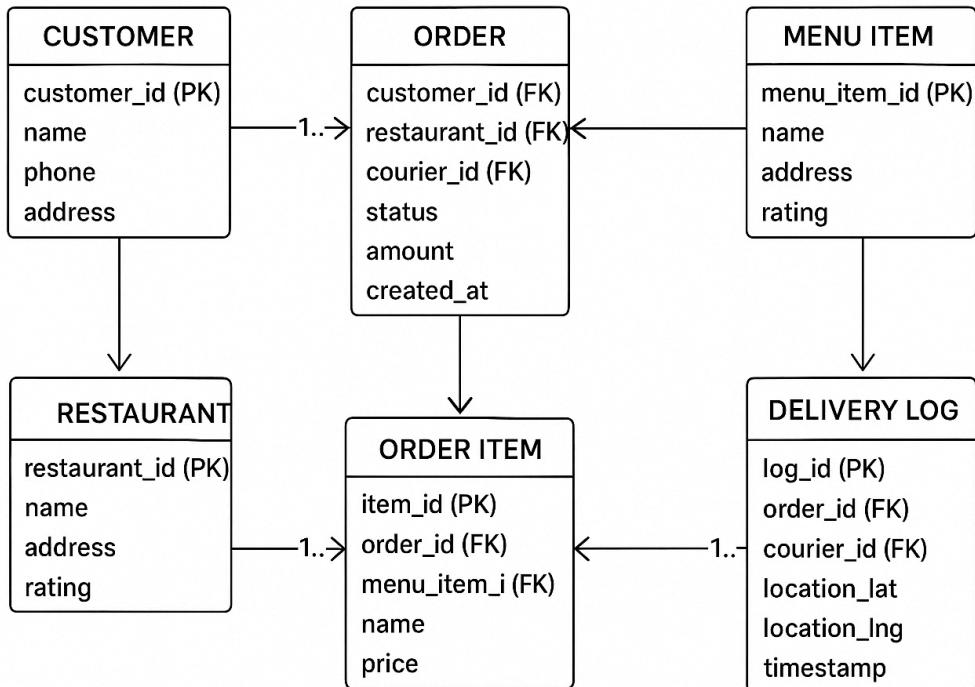


## 7) SEQUENCE DIAGRAM

Food Delivery App



## 8) ER / Data Model Diagram



## 10) Core sequence diagrams

### A. Place order

1. Customer sends POST /orders
2. API gateway → Ordering Service validates cart, reserves inventory (if needed)
3. Payment auth (optional) → Payments Service
4. Order persisted (SQL), event OrderCreated published
5. Notification & Restaurant push

### B. Dispatch & assignment

1. OrderCreated → Dispatch Service subscribes

2. Dispatch queries available couriers (Location store / redis)
3. Compute candidates + policy → assign courier (write assignment to DB)
4. Publish OrderAssigned event → Real-time gateway notifies courier & customer

## C. Live tracking

1. Courier sends periodic location updates via WebSocket
  2. Location Service writes to in-memory store (Redis) + publishes LocationUpdate
  3. Routing/ETA recalculates ETA asynchronously; updates propagated to customers via socket
- 

## 11) Dispatch design – matching algorithm

- Candidate selection: nearest N couriers within radius & available status
  - Scoring factors: ETA-to-pickup, courier load (concurrent orders), fairness score, courier acceptance rate, surge pricing incentives
  - Matching modes: greedy (fast), batch (optimize multiple orders), hybrid (time-windowed).
  - Failure handling: fallback radius expansion, requeue, manual override.
- 

## 12) Data model

- Order(id, user\_id, restaurant\_id, status, total, payment\_id, created\_at)
- OrderItem(order\_item\_id, order\_id, menu\_item\_id, qty, price)
- Menu(restaurant\_id, menu\_json, version, cached\_at)
- Courier(id, status, location(lat,lon), updated\_at, current\_load)
- Assignment(order\_id, courier\_id, assigned\_at, accepted\_at)

- RouteSnapshot(order\_id, courier\_id, eta\_pickup, eta\_delivery, polyline)

Indexes: orders(restaurant\_id, created\_at), courier(location spatial index), assignments(courier\_id, status)

---

## 13) API

- POST /v1/orders — create order (body: cart, restaurantId, paymentMethod)
- GET /v1/orders/{id} — get order status
- POST /v1/couriers/{id}/location — courier location heartbeat (or ws)
- POST /v1/dispatch/{orderId}/assign — manual assign override (admin)
- POST /v1/websocket/connect — upgrade to WS for live updates

Auth: OAuth2 / JWT tokens scoped by role (customer/restaurant/courier/admin)

---

## 14) Caching & CDN

- Menus served via CDN with short TTL + versioning on updates
  - Redis for hot reads: courier presence, slot caches, candidate caches
  - Route result caching (start,end,time) to reduce map API calls
- 

## 15) Consistency choices & tradeoffs

- Strong consistency: order state transitions & payments (ACID, SQL)
- Eventual consistency: ETA calculations, route updates, analytics

- Reason: user-facing financial flows need correctness; ETA/route are best-effort and recomputed
- 

## 16) Resiliency

- Retries with exponential backoff for external APIs
  - Circuit breakers for Maps & Payment providers
  - Idempotency keys for order placement and payment capture
  - Dead-letter queue for failed events
- 

## 17) Observability & Ops

- Structured logs (JSON), distributed traces (OpenTelemetry), metrics (Prometheus)
  - Alerts: order creation error spike, dispatch failures, map API latency, location ingestion lag
  - Dashboards for active deliveries, average ETA accuracy, courier utilization
- 

## 18) Capacity planning & rough sizing

- Target: 10k concurrent active deliveries; assume 2–5 location updates/sec per courier
- Location update QPS:  $10k * 3 = 30k$  writes/sec to ingest layer
- Use partitioned ingestion horizontally (shard by geo region)
- Kafka topics partitioned by region with 100+ partitions for throughput
- Cache tier sized for hot courier states (Redis cluster with memory to hold ~100k keys)

Back-of-envelopes included in slides.

---

## **19) SLOs & Quality targets**

- Availability: 99.9% (region); critical flows (order placement, assignment) 99.95%
  - Latency: p95 order placement < 300ms; p95 location update propagation < 1s
  - Error budget: 0.05% monthly for critical flows
- 

## **20) Maintenance & runbook**

- Blue-green deploys; schema migrations with data backfills
  - Daily integrity checks for orders vs payouts
  - On-call rotation & incident playbooks (lost courier, double-charge, payment failures)
- 

## **21) Trade-offs summary**

- SQL for strong order correctness vs NoSQL for telemetry — chosen hybrid
  - Third-party maps reduce effort but increase external dependency; mitigate with caching & fallback
  - Real-time strictness (1s) can be expensive — accept eventual consistency for analytics
-